

```

# Install the necessary libraries
!pip install --upgrade pip # Upgrade pip for latest features
!pip install --pre dlib # Install pre-built dlib if available
!pip install opencv-python-headless # Force reinstall opencv-python
!pip install -q kaggle # Install Kaggle API

🔗 Requirement already satisfied: pip in /usr/local/lib/python3.11/dist-packages (24.1.2)
Collecting pip
  Downloading pip-25.0.1-py3-none-any.whl.metadata (3.7 kB)
  Downloading pip-25.0.1-py3-none-any.whl (1.8 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 1.8/1.8 MB 32.5 MB/s eta 0:00:00
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 24.1.2
    Uninstalling pip-24.1.2:
      Successfully uninstalled pip-24.1.2
  Successfully installed pip-25.0.1
Requirement already satisfied: dlib in /usr/local/lib/python3.11/dist-packages (19.24.6)
Requirement already satisfied: opencv-python-headless in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-packages (from opencv-python-headless) (2.0.2)

# Kaggle API for installing dataset
from google.colab import files
import os

# 1. Upload kaggle.json (if you haven't already)
if not os.path.exists("/root/.kaggle/kaggle.json"):
    print("Upload your kaggle.json file")
    files.upload() # Select your kaggle.json file

# 2. Move kaggle.json to the required directory
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json
print("kaggle.json configured.")
else:
    print("kaggle.json already exists.")

# 3. Download and unzip the specific dataset
# Using the command you provided:
!kaggle datasets download davidvazquezcic/yawn-dataset -p /content/dataset --unzip
print("Dataset downloaded and unzipped to /content/dataset")

# 4. Check the structure (Important!)
# List the contents to understand the folder structure
print("\nContents of /content/dataset:")
!ls /content/dataset

🔗 Upload your kaggle.json file
  Choose Files kaggle.json
  • kaggle.json(application/json) - 69 bytes, last modified: 4/19/2025 - 100% done
  Saving kaggle.json to kaggle.json
  kaggle.json configured.
  Dataset URL: https://www.kaggle.com/datasets/davidvazquezcic/yawn-dataset
  License(s): CC-BY-NC-SA-4.0
  Dataset downloaded and unzipped to /content/dataset

  Contents of /content/dataset:
  'no yawn' yawn

# PyTorch libraries
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Dataset
from torchvision import transforms, datasets, models
from torchvision.models import ResNet18_Weights # Or other weights if using a different model

# Data handling and visualization
import numpy as np
import matplotlib.pyplot as plt
from PIL import Image
import time
import copy

# OpenCV and Dlib for real-time processing

```

```

import cv2
import dlib
from google.colab.patches import cv2_imshow # Special function for showing images in Colab

# For plotting and metrics
from sklearn.metrics import confusion_matrix
import seaborn as sns

# For saving notebook to GitHub
from google.colab import drive # Needed if saving via Drive intermediate step

print(f"PyTorch version: {torch.__version__}")
print(f"Dlib version: {dlib.__version__}")
print(f"OpenCV version: {cv2.__version__}")

# Check for GPU availability
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

↗ PyTorch version: 2.6.0+cu124
  Dlib version: 19.24.6
  OpenCV version: 4.11.0
  Using device: cuda

import os
import shutil
import random
from sklearn.model_selection import train_test_split

def split_dataset(source_dir_yawn, source_dir_no_yawn, output_dir, test_size=0.2, val_size=0.1):
    """Splits data into train, validation, and test sets.

    Args:
        source_dir_yawn: Path to the 'yawn' directory.
        source_dir_no_yawn: Path to the 'no_yawn' directory.
        output_dir: The directory where the split datasets will be saved.
        test_size: Proportion of data to include in the test split.
        val_size: Proportion of data to include in the validation split.
    """

    # Create output directories if they don't exist
    for split in ["train", "val", "test"]:
        for label in ["yawn", "no_yawn"]:
            os.makedirs(os.path.join(output_dir, split, label), exist_ok=True)

    # Splitting the yawn data
    yawn_files = [f for f in os.listdir(source_dir_yawn) if os.path.isfile(os.path.join(source_dir_yawn, f))]
    yawn_train, yawn_temp = train_test_split(yawn_files, test_size=test_size + val_size, random_state=42)
    yawn_val, yawn_test = train_test_split(yawn_temp, test_size=test_size / (test_size + val_size), random_state=42)

    # Splitting the no_yawn data
    no_yawn_files = [f for f in os.listdir(source_dir_no_yawn) if os.path.isfile(os.path.join(source_dir_no_yawn, f))]
    no_yawn_train, no_yawn_temp = train_test_split(no_yawn_files, test_size=test_size + val_size, random_state=42)
    no_yawn_val, no_yawn_test = train_test_split(no_yawn_temp, test_size=test_size / (test_size + val_size), random_state=42)

    # Copy files to respective directories
    def copy_files(files, source_dir, dest_dir):
        for file in files:
            shutil.copy(os.path.join(source_dir, file), dest_dir)

    copy_files(yawn_train, source_dir_yawn, os.path.join(output_dir, "train", "yawn"))
    copy_files(yawn_val, source_dir_yawn, os.path.join(output_dir, "val", "yawn"))
    copy_files(yawn_test, source_dir_yawn, os.path.join(output_dir, "test", "yawn"))

    copy_files(no_yawn_train, source_dir_no_yawn, os.path.join(output_dir, "train", "no_yawn"))
    copy_files(no_yawn_val, source_dir_no_yawn, os.path.join(output_dir, "val", "no_yawn"))
    copy_files(no_yawn_test, source_dir_no_yawn, os.path.join(output_dir, "test", "no_yawn"))

    # Example usage:
    source_yawn = "/content/dataset/yawn"
    source_no_yawn = "/content/dataset/no_yawn"
    output_dataset_dir = "/content/split_dataset"
    split_dataset(source_yawn, source_no_yawn, output_dataset_dir)

    # Function to visualize images from a directory
    def visualize_images(directory, num_images=5):

```

```

def visualize_images(directory, num_images=5):
    # Find image files (checking for common extensions)
    extensions = ('*.jpg', '*.jpeg', '*.png', '*.bmp', '*.gif')
    images = []
    for ext in extensions:
        images.extend([f for f in os.listdir(directory) if f.lower().endswith(ext.replace('*', '')) and os.path.isfile(os.path.join(directory, f))])

    if not images:
        print(f"No images found in {directory}")
        return

    num_images = min(num_images, len(images)) # Ensure we don't try to display more images than exist

    plt.figure(figsize=(15, 5))
    for i in range(num_images):
        img_path = os.path.join(directory, images[i])
        try:
            img = Image.open(img_path)
            plt.subplot(1, num_images, i + 1)
            plt.imshow(img)
            plt.title(f"Original:\n{images[i]}")
            plt.axis('off')
        except Exception as e:
            print(f"Error opening/displaying image {img_path}: {e}")
    plt.suptitle("Original Images Sample", fontsize=16)
    plt.tight_layout(rect=[0, 0.03, 1, 0.95]) # Adjust layout to prevent title overlap
    plt.show()

def preprocess_image(image_path):
    """
    Preprocesses an image: resizes, converts to grayscale (3 channels), normalizes.

    Args:
        image_path: The path to the image file.

    Returns:
        A preprocessed image tensor ready for a ResNet18 model or None on error.
    """
    try:
        img = Image.open(image_path).convert("RGB") # Ensure RGB
        preprocess = transforms.Compose([
            transforms.Resize((224, 224)),
            transforms.Grayscale(num_output_channels=3), # Convert to grayscale with 3 channels
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]), # Normalize
        ])
        img_tensor = preprocess(img)
        return img_tensor
    except Exception as e:
        print(f"Error processing image {image_path}: {e}")
        return None # or raise the exception

# --- Visualization setup ---
# Define directories (adjust if your path differs)
# Make sure this path exists and contains yawn images
yawn_dir = "/content/split_dataset/train/yawn"
yawn_dir = "/content/split_dataset/train/yawn" # Corrected based on previous steps
no_yawn_dir = "/content/split_dataset/train/no_yawn" # Corrected based on previous steps

# Visualize original yawn images
print(f"Visualizing original images from: {yawn_dir}")
visualize_images(yawn_dir)

# Visualize original no_yawn images
print(f"\nVisualizing original images from: {no_yawn_dir}")
visualize_images(no_yawn_dir)

# --- New code to preprocess and display one example ---

# 1. Get a sample image path (use the first image found in yawn_dir)
print("\nPreprocessing and displaying one example...")
try:
    extensions = ('*.jpg', '*.jpeg', '*.png', '*.bmp', '*.gif')
    example_image_name = None
    for ext in extensions:

```

```

potential_files = [f for f in os.listdir(yawn_dir) if f.lower().endswith(ext.replace('*', '')) and os.path.isfile(os.path
if potential_files:
    example_image_name = potential_files[0]
    break

if not example_image_name:
    raise FileNotFoundError(f"No image files found in {yawn_dir}")

example_image_path = os.path.join(yawn_dir, example_image_name)
print(f"Selected example image: {example_image_path}")

# 2. Call preprocess_image
preprocessed_tensor = preprocess_image(example_image_path)

if preprocessed_tensor is not None:
    # 3. Prepare tensor for display (denormalize and rearrange)
    # Convert tensor to numpy array
    img_np = preprocessed_tensor.numpy()

    # Transpose dimensions from (C, H, W) to (H, W, C)
    img_np = np.transpose(img_np, (1, 2, 0))

    # Denormalize the image
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    img_np = std * img_np + mean # Apply inverse transform: std * img + mean

    # Clip values to be between 0 and 1
    img_np = np.clip(img_np, 0, 1)

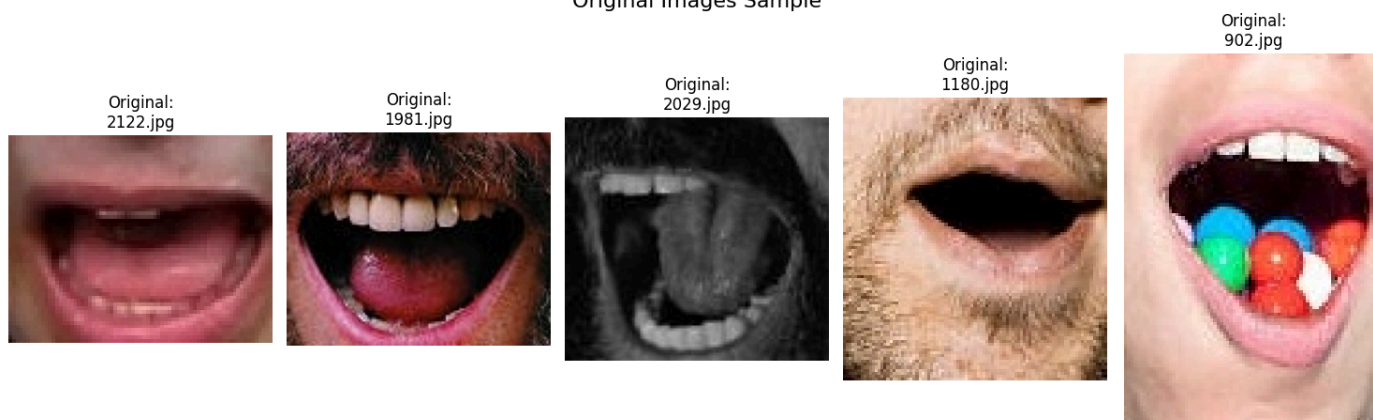
    # 4. Display using plt.imshow
    plt.figure(figsize=(6, 6))
    plt.imshow(img_np)
    # Since it's converted to grayscale with 3 identical channels, it will look grayscale
    plt.title(f"Preprocessed Example:\n{example_image_name}\n(Grayscale, Resized, Denormalized)")
    plt.axis('off')
    plt.show()
else:
    print("Preprocessing failed for the example image.")

except FileNotFoundError as fnf_error:
    print(f"Error finding example image: {fnf_error}")
except Exception as e:
    print(f"An unexpected error occurred: {e}")

```


Visualizing original images from: /content/split\_dataset/train/yawn

### Original Images Sample



Visualizing original images from: /content/split\_dataset/train/no\_yawn

### Original Images Sample




Preprocessing and displaying one example...

Selected example image: /content/split\_dataset/train/yawn/2122.jpg

Preprocessed Example:

2122.jpg

(Grayscale, Resized, Denormalized)



```

# Define the custom dataset class
# Define the custom dataset class
class YawnDataset(Dataset):
    def __init__(self, data_dir, transform=None):
        self.data_dir = data_dir
        self.transform = transform
        self.image_paths = []
        self.labels = []
        # Add the classes attribute
        self.classes = ["no_yawn", "yawn"] # Assuming "no_yawn" is 0 and "yawn" is 1

        for class_name in os.listdir(data_dir):
            class_dir = os.path.join(data_dir, class_name)
            if os.path.isdir(class_dir):
                label = 1 if class_name == "yawn" else 0
                for image_name in os.listdir(class_dir):
                    image_path = os.path.join(class_dir, image_name)
                    self.image_paths.append(image_path)
                    self.labels.append(label)

    def __len__(self):
        return len(self.image_paths)

    def __getitem__(self, idx):
        image_path = self.image_paths[idx]
        image = Image.open(image_path).convert('RGB')
        label = self.labels[idx]

        if self.transform:
            image = self.transform(image)

        return image, label

# Data Transformations
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
    'test': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
    ]),
}

# Create datasets
image_datasets = {x: YawnDataset(os.path.join("/content/split_dataset", x), data_transforms[x]) for x in ['train', 'val', 'test']}
dataloaders = {x: DataLoader(image_datasets[x], batch_size=32, shuffle=True, num_workers=4) for x in ['train', 'val', 'test']}
dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'val', 'test']}
class_names = image_datasets['train'].classes

# Load pre-trained ResNet18 model
model = models.resnet18(weights=ResNet18_Weights.DEFAULT)
num_fters = model.fc.in_features
model.fc = nn.Linear(num_fters, 2) # 2 output classes (yawn, no_yawn)
model = model.to(device)

# Loss function and optimizer
criterion = nn.CrossEntropyLoss()

```

```
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
```

```
# Training function
```

```
def train_model(model, criterion, optimizer, num_epochs=25):
    since = time.time()
    best_model_wts = copy.deepcopy(model.state_dict())
    best_acc = 0.0
```

```
    for epoch in range(num_epochs):
        print(f'Epoch {epoch}/{num_epochs - 1}')
        print('-' * 10)
```

```
        for phase in ['train', 'val']:
            if phase == 'train':
                model.train()
            else:
                model.eval()
```

```
        running_loss = 0.0
        running_corrects = 0
```

```
        for inputs, labels in dataloaders[phase]:
            inputs = inputs.to(device)
            labels = labels.to(device)
```

```
            optimizer.zero_grad()
```

```
            with torch.set_grad_enabled(phase == 'train'):
                outputs = model(inputs)
                _, preds = torch.max(outputs, 1)
                loss = criterion(outputs, labels)
```

```
            if phase == 'train':
                loss.backward()
                optimizer.step()
```

```
            running_loss += loss.item() * inputs.size(0)
            running_corrects += torch.sum(preds == labels.data)
```

```
        epoch_loss = running_loss / dataset_sizes[phase]
        epoch_acc = running_corrects.double() / dataset_sizes[phase]
```

```
        print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')
```

```
        # deep copy the model
        if phase == 'val' and epoch_acc > best_acc:
            best_acc = epoch_acc
            best_model_wts = copy.deepcopy(model.state_dict())
```

```
    print()
```

```
    time_elapsed = time.time() - since
    print(f'Training complete in {time_elapsed // 60:.0f}m {time_elapsed % 60:.0f}s')
    print(f'Best val Acc: {best_acc:4f}')
```

```
    # load best model weights
    model.load_state_dict(best_model_wts)
    return model
```

```
# Train the model
```

```
model_trained = train_model(model, criterion, optimizer, num_epochs=10) # Reduced epochs for demonstration
```

```
# Testing
```

```
def test_model(model, criterion):
    model.eval() # Set the model to evaluation mode
    running_loss = 0.0
    running_corrects = 0
    y_true = []
    y_pred = []
```

```
    with torch.no_grad():
        for inputs, labels in dataloaders['test']:
            inputs = inputs.to(device)
            labels = labels.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            loss = criterion(outputs, labels)
```

```
    running_loss += loss.item() * inputs.size(0)
    running_corrects += torch.sum(preds == labels.data)
    y_true.extend(labels.cpu().numpy())
    y_pred.extend(preds.cpu().numpy())

test_loss = running_loss / dataset_sizes['test']
test_acc = running_corrects.double() / dataset_sizes['test']
print(f'Test Loss: {test_loss:.4f} Acc: {test_acc:.4f}')

cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(8,6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=class_names, yticklabels=class_names)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("Confusion Matrix")
plt.show()

return test_acc

# Test the model
test_accuracy = test_model(model_trained, criterion)
```